



# Backend (NCCL) RoCE Fabric- High Level Design

High Performance, Lossless Ethernet Design for AI/HPC  
Workloads

Akash Patel

---

## TABLE OF CONTENTS

<b>1. Core Principles of Lossless Fabric Design.....</b>	<b>3</b>
<b>2. Fabric Topology Considerations .....</b>	<b>3</b>
<b>3. Inter GPU Communication .....</b>	<b>3</b>
NVLink/NVSwitch (intra node).....	3
RoCE.....	4
<b>4. Congestion Avoidance.....</b>	<b>4</b>
PFC, DCQCN (with ecn) & pfc watchdog .....	4
1. PFC – Priority Flow Control.....	4
2. ECN – Explicit Congestion Notification.....	5
3. DCQCN – Data Center Quantized Congestion Notification .....	5
4. PFC Watchdog .....	5
How to implement (step-by-step) .....	6
Example configs (templates; adapt to platform/versions) .....	8
<b>5. End-host (NIC).....</b>	<b>10</b>
PF, VF, and SR-IOV .....	10
Configuration.....	10
NVIDIA BlueField-3 DPUs.....	11
<b>6. Validation &amp; Acceptance Testing .....</b>	<b>12</b>
Objective .....	12
Scope .....	12
Test Categories .....	12
1. Physical & Logical Validation .....	12
2. Control Plane Validation.....	13
3. Data Plane Validation (RoCEv2 Lossless).....	13
4. Performance Benchmarking.....	13
5. Stress & Failover Testing.....	14
Documentation & Sign-Off .....	14
<b>7. Operational Best Practices .....</b>	<b>14</b>

<b><i>Appendix A- BGP DESIGN and sample configuration)</i></b> .....	<b>14</b>
<b>SAMPLE CONFIGS — Dual-Rail RoCE Backend (Rail-A shown)</b> .....	<b>15</b>
A) LEAF (Rail-A) — NCLU.....	15
B) SPINE (Rail-A) — NLCU .....	17
C) CORE (Rail-A) — NCLU .....	18
Rail-B -Configuration .....	20
<b>Quick validation (all tiers)</b> .....	<b>20</b>
<b>Design and configuration Optimization with CUMulus</b> .....	<b>20</b>

## 1. CORE PRINCIPLES OF LOSSLESS FABRIC DESIGN

**Objective:** Guarantee zero packet loss for AI training traffic (e.g., NCCL over RDMA) while maintaining high throughput and low latency.

**Key Technologies:**

- **RoCEv2** – RDMA over Converged Ethernet v2 (UDP/IP encapsulation for routing).
- **Unnumbered eBGP** – underlay routing
- **PFC (Priority Flow Control)** – Per-priority link-level pause to avoid loss.
- **ECN (Explicit Congestion Notification)** – Congestion signaling without dropping packets.

## 2. FABRIC TOPOLOGY CONSIDERATIONS

- **Clos/Leaf-Spine or Fat-Tree** for predictable oversubscription (typically 1:1).
- **Consistent hashing (ECMP)** to evenly distribute RDMA flows.
- **Dual-rail fabrics** for redundancy and increased throughput.
- **Non-blocking core** to minimize congestion points.

## 3. INTER GPU COMMUNICATION

### NVLINK/NVSWITCH (INTRA NODE)

Scope: Inside a single GPU server or chassis.

Function: Provides direct, lossless GPU-to-GPU links at hundreds of GB/s per GPU.

Topology: Fixed “islands” of 8, 16, or 32 GPUs depending on server architecture. Our solution is based on 8 GPUs per node

Design Goals:

- Maximize job placement within a single NVLink island.
- Zero packet drops/replays via vendor firmware and driver alignment.
- Maintain topology consistency across cluster.

Operational Notes:

- Monitor link counters via DCGM.
- Align CPU NUMA to GPU mapping for locality.
- Validate with NCCL intra-node tests.

## ROCE

RoCEv2 Backend Fabric (Inter-node)

Scope: Between GPU servers/chassis across the data center.

Function: Carries RDMA traffic (NCCL collectives) over Ethernet in a lossless class.

Topology: Dual-rail L3 Clos using leaf–spine–core; separate physical fabrics for compute and storage.

Design Goals:

- <300 ms failure detection/recovery using BFD with eBGP unnumbered.
- Lossless QoS: Single PFC-enabled traffic class (PCP 3), ECN tuned below PFC XOFF, DCQCN on NICs.
- Wide ECMP groups (64–128) with resilient hashing for stable NCCL flows.

Operational Notes:

- Use GSHUT for maintenance to drain traffic cleanly.
- Maintain MTU 9216 end-to-end with appropriate FEC for optics.
- Validate with `ib_write_bw`, `ib_read_lat`, and `nccl-tests` across no

## 4. CONGESTION AVOIDANCE

### PFC, DCQCN (WITH ECN) & PFC WATCHDOG

#### 1. PFC – PRIORITY FLOW CONTROL

What it is: IEEE 802.1Qbb.

- A link-level pause mechanism that applies flow control per traffic class instead of pausing the entire link.

Why for RoCE:

- RoCEv2 uses UDP over IP, so there's no built-in TCP congestion control. Without PFC, any loss in the network can cause GPU training stalls or RDMA retransmissions (very costly in HPC).

How it works:

- Switch sends a pause frame only for the lossless queue carrying RDMA traffic (e.g., PCP=3 or DSCP=26/28), while other queues (best effort, storage TCP, etc.) keep flowing.

Design note:

- Enable only on lossless class to avoid head-of-line blocking across all traffic.
- Configure symmetric PFC (both ends).
- Tune `pause_quanta` to match buffer depth.

---

## 2. ECN – EXPLICIT CONGESTION NOTIFICATION

What it is: RFC 3168

- Marks packets instead of dropping them when congestion is detected (via queue thresholds).

Why for RoCE:

- When combined with DCQCN, ECN marks trigger end-to-end rate control in RDMA flows, reducing congestion proactively.

How it works in your backend fabric:

- Switch detects queue depth over a threshold → marks ECN bits in IP header (CE codepoint).
- RDMA endpoints read these marks and slow down send rates (instead of waiting for drops).

Design note:

- Configure per-queue ECN thresholds carefully—too low and you under-utilize, too high and you risk microbursts.

Must be enabled end-to-end (every hop).

---

## 3. DCQCN – DATA CENTER QUANTIZED CONGESTION NOTIFICATION

What it is:

- NVIDIA/Mellanox implementation of congestion control for RoCEv2, using ECN marks to adjust RDMA QP send rates.

Why it matters:

- PFC alone prevents drops but doesn't stop queues from filling. DCQCN uses ECN signals to back off before PFC pause storms occur.

How it works:

- Switch ECN-marks packets.
- NIC receives ECN marks → sends Congestion Notification Packets (CNP) to sender.
- Sender NIC rate-limits that RDMA flow.

Design note:

- Enable DCQCN per lossless traffic class in NIC firmware.
- Tune min\_rate, rate\_increase, rate\_decrease parameters for your fabric RTT

---

## 4. PFC WATCHDOG

What it is:

- A safeguard that detects and breaks PFC deadlocks caused by bugs or misconfiguration.

Why it's critical:

- PFC deadlocks can bring an entire fabric class to a halt—watchdog forces recovery.

How it works:

- Monitors paused queues for a timeout (e.g., >500ms).
- If timeout hit → flushes queue, drops frames, logs event.

Design note:

- Enable on all switches in lossless queues.
- Tune timeout so it doesn't trigger on normal pauses, only pathological cases.

#### How They Work Together in Your RoCE Backend

Function	Scope	Prevents	Side Effect if Misused
PFC	Link-level	Packet loss in lossless queues	Head-of-line blocking, deadlocks
ECN	End-to-end	Long-term queue buildup	Under-utilization if threshold too low
DCQCN	Endpoint rate control	PFC storms, persistent congestion	Lower throughput if too aggressive
PFC Watchdog	Switch queue safety	Fabric lock-up from PFC deadlock	Possible packet loss during flush

## HOW TO IMPLEMENT (STEP-BY-STEP)

### 1) Classify traffic (one lossless class only)

- Mark **RoCEv2** as **DSCP 26 (AF31)** → **PCP 3** (example).
- Keep storage/control/management **lossy** (no PFC).

### 2) MTU & hashing

- Set MTU 9000/9216 end-to-end (hosts + switches).
- Ensure ECMP/LAG hashing has good entropy (5-tuple including UDP).

### 3) Buffers & PFC thresholds (XOFF/XON)

#### Goal

Size headroom so a receiver (switch/NIC) can absorb all bytes “in flight” after it sends PFC pause—without loss—then pick sane XOFF/XON marks for the lossless priority queue.

#### Headroom formula

Headroom  $\geq$  (WireRate (bits/s) × ReactionWindow (s)) / 8 + WorstCaseBurst (bytes)

- ReactionWindow  $\approx$  25–35  $\mu$ s (pause frame TX + propagation + upstream stop + drain start)
- WorstCaseBurst: 64–128 KB (multi-sender fan-in, hash collisions, pipeline flush, etc.)

#### Mapping Headroom → XOFF/XON

- Pick XOFF a bit below headroom to account for jitter/measurement error.
- Pick XON so there's enough drain (XOFF – XON) to clear the queue before senders resume.

#### Rule of thumb

- XOFF  $\approx$  70–85% of Headroom
- XON  $\approx$  35–50% of Headroom
- Drain time  $\approx$  (XOFF – XON) / LineRate (bytes/s)

#### examples

##### 100 G

- Headroom target:  $\approx$  450–500 KB
- Suggested marks: XOFF = 320–384 KB, XON  $\approx$  192 KB

- Drain  $\approx 130\text{--}190\text{ KB} \rightarrow \sim 10\text{--}15\text{ }\mu\text{s}$  at 100 G)
- 200 G
- Headroom target:  $\approx 900\text{--}1000\text{ KB}$
  - Suggested marks: XOFF =  $700\text{--}850\text{ KB}$ , XON  $\approx 400\text{ K}$
  - Drain  $\approx 300\text{--}450\text{ KB} \rightarrow \sim 12\text{--}18\text{ }\mu\text{s}$  at 200 G)

#### Notes / definitions

- PFC TX: Receiver (switch/NIC) detects its egress queue for the lossless priority approaching XOFF, then transmits a PFC pause frame upstream.
- Drain time: Time (and bytes) between XOFF and XON while the queue empties.
- Upstream stop: Time for senders to process pause, flush TX pipelines, and cease new frames.
- Worst-case burst: Additional data that can still arrive after pause due to multiple senders, pipeline, and scheduling artifacts (assume  $64\text{--}128\text{ KB}$ ).

#### Practical tips

- Size per priority (not global buffer).
- Re-verify after changing MTU, number of senders, or ECMP fan-in.
- If using dual-rail fabrics, validate each rail independently.
- Watch tail latency: too-high XOFF/XON can amplify PFC storms; too-low risks drops.

#### 4) ECN/DCQCN first, PFC second

- Enable ECN on the lossless queue to mark before PFC triggers.
- Start  $K_{\min} \sim 10\text{--}20\%$  below XOFF;  $K_{\max} \sim 2\text{--}3 \times K_{\min}$  but  $<$  headroom
- 100G example:  $K_{\min} 64\text{ KB}$ ,  $K_{\max} 192\text{ KB}$
- 200G example:  $K_{\min} 128\text{ KB}$ ,  $K_{\max} 320\text{ KB}$
- Enable DCQCN on hosts (ConnectX): conservative target rates so queues stay shallow.

#### 5) Enable PFC on links carrying RoCE

- Turn **PFC on only for that one priority** (e.g., PCP 3). Everything else remains lossy.

#### 6) Enable a PFC watchdog (deadlock breaker)

- Detect continuous pause on the lossless class and **temporarily drop or make that class lossy** to drain.
- Start with **detect 200 ms, restore 2 s, action: drop** (or “lossy”).

#### 7) Host (NIC) alignment

- Map DSCP 26  $\rightarrow$  PCP 3.
- Enable RoCEv2 + DCQCN; PFC only on the RoCE priority.
- Verify jumbo MTU and GID/IP config.

#### 8) Validate

- Run incast/NCCL tests at  $\sim 90\text{--}95\%$  link rate.
  - Watch per-priority PFC rx/tx, ECN CE marks, watchdog events, RDMA retrans/timeouts.
  - Good sign: steady ECN, near-zero PFC, watchdog never triggers, zero RDMA drops.
-



---

## EXAMPLE CONFIGS (TEMPLATES; ADAPT TO PLATFORM/VERSIONS)

### Cisco NX-OS

```
class-map type qos match-any ROCE-CLASS
  match dscp 26

policy-map type qos ROCE-QOS
  class ROCE-CLASS
    set qos-group 3

policy-map type network-qos ROCE-NQ
  class type network-qos class-default
    mtu 9216
  class type network-qos qos-group 3
    pause no-drop
    pfc-cos 3
    ecn
    ecn mark-probability 100 min-threshold 64K max-threshold 192K
system qos
  service-policy type qos input ROCE-QOS
  service-policy type network-qos ROCE-NQ
# Set XOFF/XON for the no-drop class (ASIC/queue CLI varies by platform)
```

NVIDIA Spectrum (SONiC example for watchdog; Cumulus shown for QoS/PFC)

### Cumulus/NCLU (QoS/PFC/ECN):

```
net add dcb pfc priority 3 enable
net add qos map dscp-to-tc 26:3
net add qos ecn enable
net add qos ecn queue 3 min-threshold 64k max-threshold 192k mark-probability 100
net add interface swp1-48 mtu 9216
net pending; net commit
```

### SONiC (PFC watchdog):

```
pfcwd start --action drop --restoration-time 2000
pfcwd interval 100
pfcwd start --priority 3 --detection-time 200
```

### Arista EOS

```
qos map dscp-to-traffic-class 26 3
priority-flow-control enable
priority-flow-control priority 3 enable
queue-monitor length ecn 3 min-threshold 64 kilobytes max-threshold 192 kilobytes
```

### # Platform-specific queue headroom / xoff/xon commands here

```
Host (Linux / ConnectX)
# Map DSCP 26 → PCP 3
dcb app add dev <ifc> dscp 26 prio 3
```

```
# Enable RoCEv2 + DCQCN (exact knobs vary by NIC fw)
mlxconfig -d <pci> set ROCE_ACCL=1
mlxconfig -d <pci> set DCBX_PFC_MODE=1 ETS_PFC_ENABLE=0x08 # PFC on prio 3
ip link set dev <ifc> mtu 9216
```

## 5. END-HOST (NIC)

### PF, VF, AND SR-IOV

#### 1. PF – Physical Function

- The real PCIe device, as seen by the OS.
  - Has full hardware control: can configure QoS, PFC, DCQCN, SR-IOV settings, VLAN mappings, firmware features.
  - Each physical NIC port has one PF.
  - In Mellanox/NVIDIA terms:
    - The PF is where you run `mlxconfig`, enable PFC, set DSCP→PCP mappings, etc.
    - Whatever you configure here can be pushed down to VFs.
- 

#### 2. VF – Virtual Function

- A lightweight “slice” of the physical NIC created by SR-IOV.
  - Appears as an independent PCIe device to a VM, container, or tenant OS.
  - Has its own MAC, queue pairs, interrupts, but cannot change deep hardware policy (that’s the PF’s job).
  - In RoCE terms:
    - A VF can do RDMA (with QoS and congestion control inherited from the PF).
    - You can map different VFs to different VLANs/priorities if the PF is set up to allow it.
- 

#### 3. SR-IOV – Single Root I/O Virtualization

- PCIe standard that allows one PF to present multiple VFs to the system.
- Lets each VF bypass the host kernel’s slow path and talk directly to NIC hardware (hardware offload).
- Benefits in AI/ML and HPC fabrics:
- Low latency (no software switching in the hypervisor).
- High throughput (direct DMA from VF to NIC queues).
- Hardware isolation per VF (each VM/container sees “its own” NIC).

### CONFIGURATION

**Two clear classes** at the NIC:

- **Lossless class:** RDMA QPs for RoCEv2 traffic → Priority 3 (example).
- **Lossy class:** Everything else (TCP/IP, management).

Example of **NVIDIA/Mellanox ConnectX-4 Lx and newer** (ConnectX-4/5/6/6 Dx/7)

#### 1. Classify RDMA traffic

Map RoCEv2 DSCP (26) to lossless PCP (3)

- `dcb app add dev eth0 dscp 26 prio 3`

2. Map other traffic to lossy priorities

Example: map management and TCP/IP to prio 0 (lossy)

- `dcb app add dev eth0 dscp 0 prio 0`

3. Enable PFC only on lossless prio

- `mlxconfig -d <pci_bus> set DCBX_PFC_MODE=1 ETS_PFC_ENABLE=0x08 # bit mask for prio 3`

4. Confirm mappings

- `dcb app show dev eth0`

## NVIDIA BLUEFIELD-3 DPUS

Nvidia Bluefield 3 (BF3) use essentially the same RoCEv2, DCQCN, and per-priority PFC configuration model as the ConnectX NICs, because the DPU's embedded networking pipeline is ConnectX-7-class under the hood. Here is the comparison

Feature / Step	ConnectX-5/6/7 NIC	BlueField-3 DPU
<b>Where you run commands</b>	Host OS (x86/Arm server)	On the DPU Arm cores (Ubuntu/DPU-OS) or via DOCA
<b>Firmware tool</b>	mlxconfig (part of MFT package)	mlxconfig on DPU OS (same syntax)
<b>Verify RoCE support</b>	<code>`mlxconfig -d &lt;pci&gt; query</code>	<code>grep ROCE`</code>
<b>Enable RoCEv2</b>	<code>mlxconfig -d &lt;pci&gt; set ROCE_ACCL=1 ROCE_IPV6_ONLY=0</code>	Same
<b>Enable DCQCN</b>	<code>mlxconfig -d &lt;pci&gt; set CNP_PACING_MODE=1</code>	Same
<b>Map DSCP→PCP</b>	On host: <code>dcb app add dev &lt;ifc&gt; dscp 26 prio 3</code>	On DPU OS: same <code>dcb</code> or <code>tc</code> command for PFs; for VFs, map in PF and propagate to VF
<b>Enable PFC (lossless class only)</b>	<code>mlxconfig -d &lt;pci&gt; set DCBX_PFC_MODE=1 ETS_PFC_ENABLE=0x08 (bitmask for PCP 3)</code>	Same on DPU OS; applies to PF; ensure VF inherits or override in VF config
<b>ECN thresholds (switch side)</b>	Kmin ≈ 64 KB, Kmax ≈ 192 KB for 100 G (below XOFF)	Same — BlueField NIC pipeline behaves like ConnectX-7; thresholds set on switch, not DPU

<b>PFC watchdog</b>	On switch OS (e.g., SONiC/Arista/Cisco); not NIC-local	Same — controlled in the switch, not the DPU
<b>SR-IOV mapping</b>	PF owns PFC/DCQCN config; VFs inherit	PF config on DPU controls VF behavior; can also isolate via DOCA Flow
<b>Telemetry</b>	ethtool -S, mlxstat on host	ethtool -S, mlxstat on DPU OS; DOCA Telemetry for per-priority counters
<b>Special notes</b>	Runs entirely in host NIC	DPU may terminate control/management traffic locally; ensure only RDMA data path is in lossless class

## 6. VALIDATION & ACCEPTANCE TESTING

### OBJECTIVE

Verify that the backend RoCE fabric (NCCL traffic domain) meets design intent for lossless transport, resiliency, and performance under operational conditions before being placed into production.

### SCOPE

Covers dual-rail backend fabric interconnecting GPU nodes across 98 pods / 392 SU. Tests validate:

- Correct physical & logical connectivity (port, VLAN/VRF assignments).
- RoCEv2 lossless behavior (PFC, ECN, DCQCN tuning).
- BGP underlay stability and fast reconvergence.
- NCCL collective communication performance and path diversity.

### TEST CATEGORIES

#### 1. PHYSICAL & LOGICAL VALIDATION

Checkpoint	Method	Acceptance Criteria
Port link-up / MTU	ethtool, show interfaces	All fabric links operational at 400G; MTU=9216

VLAN / VRF mapping	net show bridge vlan, ip vrf show	Backend VRF correctly applied to all RoCE interfaces
Optics & FEC	ethtool -m, NOS optics diag	Optics match design spec; FEC enabled as per vendor guidance

## 2. CONTROL PLANE VALIDATION

Checkpoint	Method	Acceptance Criteria
eBGP unnumbered peering	net show bgp summary	All leaf–spine–core adjacencies up
BFD timers	net show bfd	Sub-300 ms failure detection
ECMP path count	net show bgp ipv4 unicast	ECMP width matches ASIC limit (≥64)
Graceful maintenance	Simulated GSHUT	Traffic drains cleanly with no packet loss

## 3. DATA PLANE VALIDATION (ROCEV2 LOSSLESS)

Checkpoint	Method	Acceptance Criteria
PFC pause frame counters	ethtool -S or mlxreg	No sustained pause storm; counters stable. Mlxreg for advance troubleshooting
ECN marking	ethtool -S, telemetry collector	ECN marks in acceptable range under load
DCQCN behavior	Flow replay under congestion	Congestion window ramps down/up smoothly

## 4. PERFORMANCE BENCHMARKING

Tool	Purpose	Acceptance Criteria
ib_write_bw / ib_write_bw -d mlx5_0 -F / ib_read_lat	RDMA microbenchmarks	<p>Bandwidth ≥ link-rate efficiency; latency &lt; design target</p> <p>** for ConnectX-5(6/7) install perftest using apt-get (ubuntu/Debian) or yum (RHEL/Cent-OS) before writing this</p> <p>If BlueField3, then you can directly SSH into it</p>

nccl-tests (all_reduce, all_gather)	Application-level NCCL validation	Throughput within 5% of lab reference; no retries triggered
Multipath diversity test	Parallel NCCL jobs	Flows spread evenly across ECMP paths

## 5. STRESS & FAILOVER TESTING

Scenario	Method	Acceptance Criteria
Link failure	Disable leaf–spine link	No job failure; reconvergence < 300 ms
Node reboot	Restart GPU node	NCCL jobs recover without hang
Multi-link failure (within ECMP width)	Disable multiple spines	Bandwidth reduction proportional; no packet loss in remaining paths

## DOCUMENTATION & SIGN-OFF

- All test results recorded in Validation Report with:
  - Test date/time
  - Node pair / fabric rail tested
  - Tool output logs
  - Pass/Fail summary

## 7. OPERATIONAL BEST PRACTICES

- Keep firmware and switch OS updated for latest RDMA/PFC fixes.
- Monitor per-priority buffer utilization, ECN marks, and pause events.
- Avoid enabling PFC globally—scope it tightly to RDMA priorities.
- If congestion hotspots appear, evaluate load balancing (e.g., adaptive routing).

## APPENDIX A- BGP DESIGN AND SAMPLE CONFIGURATION)

Design intention (why)	Implementation on Spectrum/Cumulus
------------------------	------------------------------------

Fast failure detection < 200–300 ms	BFD on every eBGP peer: tx/rx 50–100 ms, multiplier 3
Blast radius containment	Dual rails, separate ASNs per rail; no interconnect between rails
L3 CLOS with wide load-balancing	eBGP unnumbered on all leaf↔spine↔core links; ECMP 64; multipath-relax
Minimal flow churn on path changes	Enable resilient/consistent ECMP hashing (if exposed); include UDP L4 ports in hash
Preserve dataplane during control restarts	BGP Graceful Restart (GR) on all tiers (Leaf/Spine/Core)
Immediate removal of failed paths	Fast external fallover; drop eBGP adjacency on link-down
Maintenance without packet loss	GSHUT (graceful-shutdown) before interface or process work
Lossless only for RoCE	One lossless class (PCP 3) mapped from DSCP (26/40/44 → 3); PFC only on that class
Keep queues short, avoid PFC churn	ECN on TC3 with Kmin/Kmax below PFC XOFF; hosts/DPUs run DCQCN
Jumbo & physical hygiene	MTU 9216 end-to-end; optics-appropriate FEC; watch micro-loss
Simple, scoped routing	Redistribute only SVIs (leaf) and loopbacks; no non-underlay routes
Observability	Stream per-priority PFC/ECN + queue stats; alarms for BFD down, ECMP shrink, prefix churn

## SAMPLE CONFIGS — DUAL-RAIL ROCE BACKEND (RAIL-A SHOWN)

Mirror the same pattern for **Rail-B** with its **own ASN**, VLAN/SVI space, and physical links.

Backend VRF: vrf-backend-nccl. Lossless class: **TC/PCP 3**. DSCP for RoCE: **26** (or **40/44** if you prefer). This is using Nvidia Cumulus (NCLU) for illustration purpose, but similar configuration can be applied to FRR or other vendor switching solution

### A) LEAF (RAIL-A) — NCLU

#### Assumptions

- Leaf-A ASN 65011; loopback 10.0.11.1/32
- Host RDMA VLAN/SVI (this SU/rail): VLAN 1005, 10.100.5.1/26, MTU 9216
- Uplinks to spines: swp33–swp34 (routed)



- Host ports: swp1–swp32(access VLAN 1005)

## # System & Loopback

```
net add loopback lo ip address 10.0.11.1/32
```

```
# VRF & SVI (backend RDMA for this SU/rail)
```

```
net add vlan 1005
net add interface vlan1005 ip address 10.100.5.1/26
net add interface vlan1005
net add interface vlan1005 mtu 9216
```

```
# Host access ports
```

```
net add bridge bridge vids 1005
net add bridge bridge ports swp1-32
net add interface swp1-32 bridge access 1005
net add interface swp1-32mtu 9216
```

```
# Uplinks to Spines (routed, unnumbered eBGP)
```

```
net add interface swp33-64 mtu 9216
```

```
# QoS / Lossless isolation (RoCE only)
```

```
net add dcb pfc priority 3 enable
net add qos map dscp-to-tc 26:3 40:3 44:3
net add qos ecn enable
net add qos ecn queue 3 min-threshold 64k max-threshold 192k mark-probability
100
```

```
# BGP core: eBGP unnumbered + BFD + ECMP + GR + GSHUT (ready)
```

```
net add bgp autonomous-system 65011
net add bgp router-id 10.0.11.1
net add bgp bestpath as-path multipath-relax
net add bgp graceful-restart # if control plan
crash/restart , the data plane continue to work (FIB)
net add bgp graceful-shutdown # enable/disable during drains ,
used during maintenance window
```

```
net add bgp neighbor SPINES peer-group
net add bgp neighbor SPINES remote-as external
net add bgp neighbor SPINES timers 60 180
```

```
for i in {33..64}; do
    net add bgp neighbor swp$i interface peer-group SPINES
    net add bgp neighbor swp$i bfd
done
```

```
# Address-family
```

```
net add bgp address-family ipv4 unicast maximum-paths 64
net add bgp address-family ipv4 unicast redistribute connected
```

```
# Commit
```

```
net commit
```

```
Operational drain (GSHUT)
```

```
# Start graceful shutdown to drain traffic
```

```
net add bgp graceful-shutdown
```

```
net commit
```

```
# After maintenance:
```

```
net del bgp graceful-shutdown
```

```
net commit
```

---

## B) SPINE (RAIL-A) — NLCU

Assumptions

- Spine ASN 65100; loopback 10.0.100.1/32
- Downlinks to leaves: swp1-swp32 (routed)
- Uplinks to core: swp49-swp64 (routed)

```
# System & Loopback
```

```
net add loopback lo ip address 10.0.100.1/32
```

```
# Links
```

```
net add interface swp1-64 mtu 9216
```

```
# QoS policy (mirror lossless class across fabric)
```

```
net add dcb pfc priority 3 enable
```

```
net add qos map dscp-to-tc 26:3 40:3 44:3
```

```
net add qos ecn enable
```

```
net add qos ecn queue 3 min-threshold 64k max-threshold 192k mark-probability 100
```

```
# BGP core
```

```
net add bgp autonomous-system 65100
```

```
net add bgp router-id 10.0.100.1
```

```
net add bgp bestpath as-path multipath-relax
```

```
net add bgp graceful-restart
```

```
net add bgp graceful-shutdown
```

```

# Peer-groups
net add bgp neighbor LEAFS peer-group
net add bgp neighbor LEAFS remote-as external
net add bgp neighbor LEAFS timers 60 180

net add bgp neighbor CORES peer-group
net add bgp neighbor CORES remote-as external
net add bgp neighbor CORES timers 60 180

# Attach downlinks to LEAFS + BFD
for i in {1..32}; do
    net add bgp neighbor swp$i interface peer-group LEAFS
    net add bgp neighbor swp$i bfd
done

# Attach uplinks to CORES + BFD
for i in {33..64}; do
    net add bgp neighbor swp$i interface peer-group CORES
    net add bgp neighbor swp$i bfd
done

# Address-family
net add bgp address-family ipv4 unicast maximum-paths 64
net add bgp address-family ipv4 unicast redistribute connected

net commit
Operational drain (GSHUT)
net add bgp graceful-shutdown
net commit

# remove after:
net del bgp graceful-shutdown
net commit

```

---

## C) CORE (RAIL-A) — NCLU

### Assumptions

- Core ASN 65200; loopback 10.0.200.1/32
- Downlinks to spines: swp1-swp64 (routed)
- Core carries underlay only (no SVIs); optional service loopbacks (telemetry, etc.) are fine to advertise
- Same QoS policy (TC3 lossless) to keep PFC semantics consistent across the fabric

### # System & Loopback

```
net add loopback lo ip address 10.0.200.1/32

# Downlinks to Spines with Breakout each 800G port into 2x400G
for p in {1..64}; do
    net add interface swp$p breakout 2x400G
done

net pending
net commit

# Set MTU for each breakout child port
for p in {1..64}; do
    for s in 0 1; do
        net add interface swp${p}s${s} mtu 9216
    done
done

net add dcb pfc priority 3 enable
net add qos map dscp-to-tc 26:3 40:3 44:3
net add qos ecn enable
net add qos ecn queue 3 min-threshold 64k max-threshold 192k mark-probability
100

# BGP core
net add bgp autonomous-system 65200
net add bgp router-id 10.0.200.1
net add bgp bestpath as-path multipath-relax
net add bgp graceful-restart
net add bgp graceful-shutdown

# Peer group for spines
net add bgp neighbor SPINES peer-group
net add bgp neighbor SPINES remote-as external
net add bgp neighbor SPINES timers 60 180

# Attach all spine-facing interfaces + BFD
for p in {1..64}; do
    for s in 0 1; do
        net add bgp neighbor swp${p}s${s} interface peer-group SPINES
        net add bgp neighbor swp${p}s${s} bfd
    done
done

# Address-family
net add bgp address-family ipv4 unicast maximum-paths 64
# Only advertise loopbacks (and any explicit service loopbacks)
net add bgp address-family ipv4 unicast redistribute connected

net commit
Operational drain (GSHUT)
net add bgp graceful-shutdown
net commit
# remove after:
net del bgp graceful-shutdown
net commit
```

---

## RAIL-B -CONFIGURATION

- Rail-B: duplicate each block with its own ASNs (e.g., Leaf-B 65012, Spine-B 65101, Core-B 65201), its own VLAN/SVI space (e.g., 10.100.6.0/26 for the Leaf-B SU), and separate physical links.
- No inter-rail connectivity: rails are independent failure domains. Any cross-rail services (rare) should live outside the RoCE underlay.

**\*\* [Same design and configuration with different Loopback & /26 subnets \(NVM-of RoCE\) would be applicable for Storage Fabric](#)**

## QUICK VALIDATION (ALL TIERS)

```
# BGP/BFD health
net show bgp summary
net show bfd peers

# ECMP width / routes
net show route 10.100.5.0/26
net show route summary

# QoS / lossless signals
sudo ethtool -S swp1 | egrep -i 'pfc|prio|pause'
sudo cl-qos-show
```

## DESIGN AND CONFIGURATION OPTIMIZATION WITH CUMULUS

### BGP unnumbered

This example leverages BGP Unnumbered to optimize the configuration and save thousands of /31 IP addresses.

- No IPv4 / IPv6 global address needed on the interface
- The switch interface only uses IPv6 link-local addresses (FE80::/10), which are automatically created for every L3 interface.
- These link-local addresses exist without you explicitly assigning anything.
- When you do bfd on that neighbor, it also uses IPv6 link-local over the same interface.

### BGP external

Normal BGP requires you to define the remote AS number as an example `net add bgp neighbor swp33 remote-as 65102`

- However, in our sample configuration, remote-as external tells FRR:
- “This neighbor’s AS number is different from local AS. Learn it dynamically during session establishment.”
- This is valid for eBGP sessions only, because eBGP expects AS numbers to differ.
- The local router’s ASN is already set (e.g., `net add bgp autonomous-system 65101`).
- On TCP session open, FRR inspects the remote’s BGP Open message, reads the ASN from the packet, and uses it without you hardcoding it